

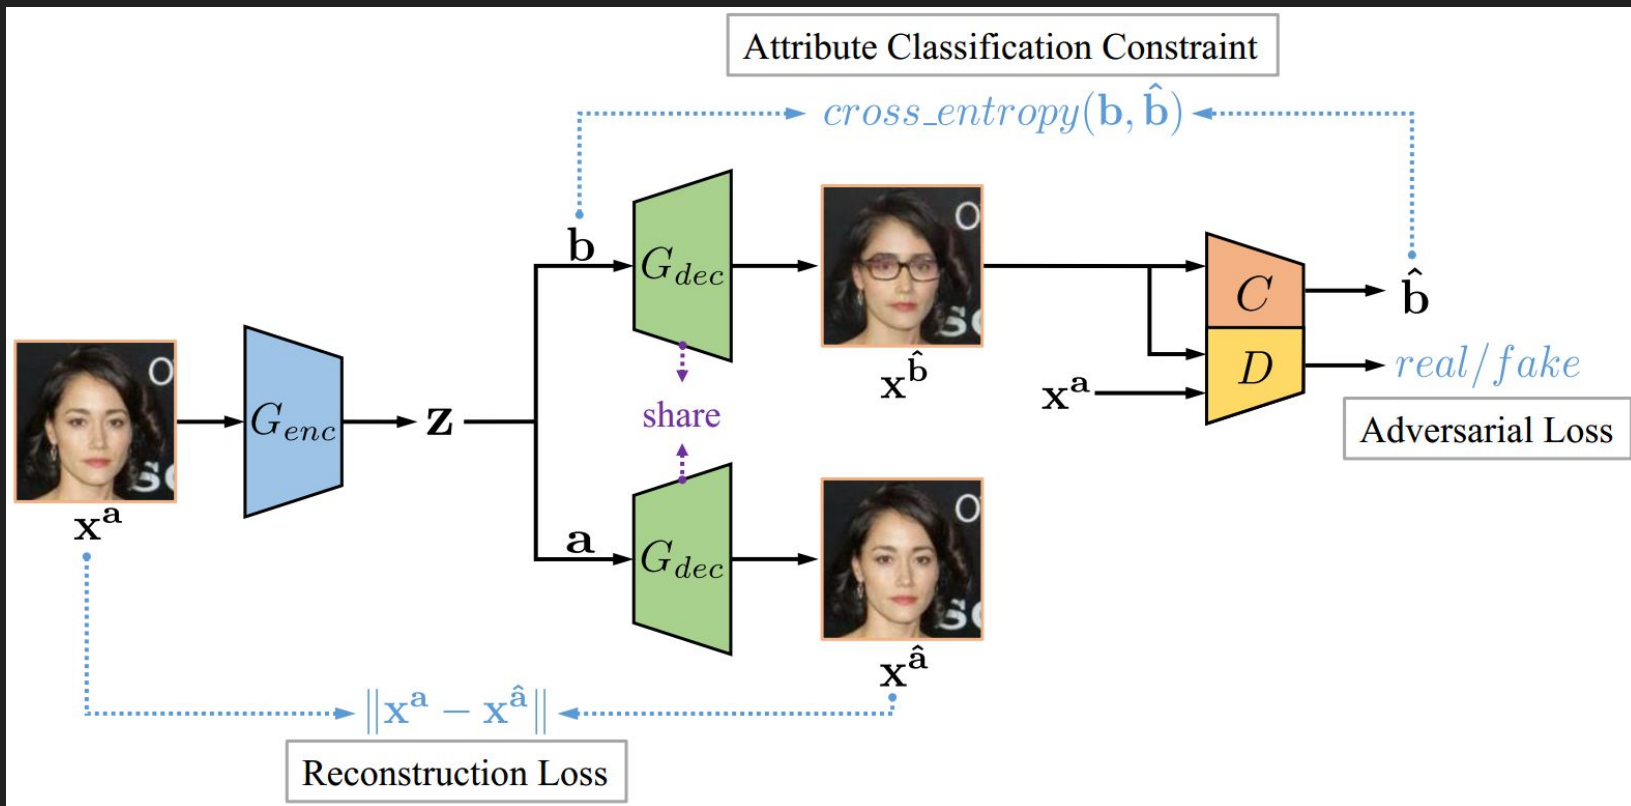
Facial Attribute Editing using Neural Networks

Generative Adversarial Networks

Contents

- Intro to Model and How it works
- Training process
- Loss Metrics
- Graphs
- Generated Samples

Model



Intro

- *Networks:*
 - G-enc
 - Encodes the Image to Latent space(Z)
 - G-dec
 - Decodes the Image to original Image Conditioned on specified attributes
 - C & D
 - Attribute Classification constraint works on logits captured at end Layer
 - Real/Fake Detector
- *Loss functions*
 - Reconstruction Loss which preserves the original attributes
 - Adversarial Loss for Real like Sample (WGAN)
 - Cross entropy for Attribute Classification

Training Process

- Load the batch of images (Image[?,128,128,3]+Attribute[?,len(Desired_Attributes)])
- Generate Valid Random Attribute vectors
- Main Three Steps

Xa: Original Image


- Generate samples for Original and Random Attribute
 - $\text{Generated_image}_{\text{original}} = \text{Gdec}\{\text{Genc}(\text{Xa}), \text{original Attribute vector}\}$
 - $\text{Generated_image}_{\text{random}} = \text{Gdec}\{\text{Genc}(\text{Xa}), \text{random Attribute vector}\}$
- Discriminate samples for Original and Random Attribute
 - $\text{Xa_logits}_{\text{gan}}, \text{Xa_logits}_{\text{attr}} = \text{D}(\text{Xa}), \text{C}(\text{Xa})$
 - $\text{Xb_logits}_{\text{gan}}, \text{Xb_logits}_{\text{attr}} = \text{D}(\text{Generated_image}_{\text{random}}), \text{C}(\text{Generated_image}_{\text{random}})$
- Calculate Reconstruction loss
- Define Loss functions Three steps
 - $\text{D-loss} = (\text{Xa_logits}_{\text{gan}} - \text{Xb_logits}_{\text{gan}}) + 10.0 * \text{GP} + \text{Xa_loss}_{\text{attr}}$ (More in Next Slide)
 - $\text{G-loss} = \text{Avg}(\text{Xb_logits}_{\text{gan}}) + 10.0 * \text{Xb_loss}_{\text{attr}} + 100 * \text{Reconstruction_loss}$ (More in next Slide)
- Define Optimizers (Adam)
 - $\text{D-step} = \text{tf.train.AdamOptimizer}(\text{lr}, \text{beta1}=0.5). \text{minimize}(\text{d_loss}, \text{var_list}=\text{d_var})$
 - $\text{G-step} = \text{tf.train.AdamOptimizer}(\text{lr}, \text{beta1}=0.5). \text{minimize}(\text{g_loss}, \text{var_list}=\text{g_var}, \text{global_step}=\text{global_step})$
- Run for epoch in n_epochs:
 - Run for batch in n_batches:
 - Run D-step : $\text{G_step} = 5 : 1$

Loss Metrics


- WGAN-GP
 - $L_D^{\text{WGAN}} = E[D(X_a)] - E[D(\text{Generated_image}_{\text{random}})]$
 - $L_G^{\text{WGAN}} = E[D(\text{Generated_image}_{\text{random}})]$
 - Gradient Penalty = $\lambda * (\|\nabla_x (D(x))\|_2 - 1)^2$
 - Total_loss = $L_D^{\text{WGAN}} + L_G^{\text{WGAN}} + \text{Gradient Penalty}$
- Xa_loss_{attr} : `tf.losses.sigmoid_cross_entropy(a, Xa_logit_attr)`
- Xb_loss_{attr} : `tf.losses.sigmoid_cross_entropy(a, Xa_logit_attr)`

Graphs


Xa_loss_Reconstruction




Xa_loss_att




D_loss_gan



Xb_loss_Gan



Xb_loss_Attribute



Generated Samples



Original

Original

Bangs

Bald

Black Hair

Blond hair

BrownHair

Bushy
EyeBrows

Specs

Gender

Mouth open
(Slight)

Mustache

Beard

Pale Skin

Age

End

